

MechanicaFluidorum

A Verifier-in-the-Loop Program Motivated by 3D Navier–Stokes Regularity

Dual-Scale Geometry, the Symmetric-Square Lock, and
the Truncated Viscous Katz–Pavlović Dyadic Shell Model

*Scientific Progress Report — Intermediate Results,
Revised After the External Audit of 13 August 2026*

Xavier Callens

SocrateAI Scientific Programme

SocrateAI-Scientific-MechanicaFluidorum

13 August 2026

Repository state: commit `a35cb92+` | All verification gates PASS (exit 0)

Abstract

This report documents the state of a staged, verifier-in-the-loop research programme on the global regularity of the three-dimensional incompressible Navier–Stokes equations (NSE), organised around *Hypothesis U* — a uniform-in-cutoff enstrophy bound for the frequency-truncated system. The programme’s distinguishing feature is not a claimed proof but a *method*: every assertion carries an explicit epistemic tier, mechanically enforced by a two-gate verification pipeline (exact rational arithmetic; Lean 4 kernel compilation with a pinned axiom footprint). We report 72 theorems kernel-compiled across seven Lean files — of which 65 carry Tier A claims, the remaining 7 having been *demoted* to Tier C by the audit below — six exact-arithmetic certification harnesses, and, equally importantly, a documented record of what has *failed*: two abandoned numerical designs, five caught errata (one in a formula sourced from the literature, one in the programme’s own design memo, one a broken commit, one a false positive in the verification gate itself, and one an unpinned dependency that would have verified the programme against the wrong prover library), four analytical tracks that remain correctly blocked on undefined objects, and — the largest single failure in the programme’s history — **an external expert audit that killed this report’s own central framing claim (§2). The retraction is stated in full and accepted without qualification**: the formalisation, with a one-dimensional index cutoff, does not restate the unreduced three-dimensional problem; it describes a dyadic shell hierarchy. The programme has been re-targeted accordingly, and its formal object is now the truncated viscous Katz–Pavlović dyadic shell model. **No claim is made that Hypothesis U is true, in any form, nor that the Navier–Stokes problem is resolved.** The Millennium Prize problem remains open and is not addressed by anything proven here.

Contents

1	Context Management: Reading This Document and This Repository	3
1.1	The governing documents and their precedence	3
1.2	The epistemic charter	3
1.3	The three-tier gating system	4
1.4	The honesty clause and the obstruction ledger	4
1.5	Reproducing the verification state	5

2	The Problem, Stated Honestly	5
2.1	Classical setting	5
2.2	The regularised family and Hypothesis U	6
3	Dual-Scale Geometry and T-Duality	7
3.1	Physical motivation (Tier C) and its strict boundary	7
3.2	The effective radius: what is actually proven (Tier A)	7
3.3	The Symmetric-Square Lock (Tier A)	8
4	The “Holographic Algebra” Organising Concept	9
4.1	The picture	9
4.2	Why this picture would matter, <i>if</i> it could be substantiated	9
5	The Dyadic Shell Model: Established Results	9
5.1	Energy flux telescoping and monotonicity (Tier A)	10
5.2	The exact enstrophy-production identity (Tier A)	10
5.3	A local production bound (Tier A)	11
6	Statement-Level Formalisation, the Pivot, and the Demotion	11
6.1	The shell-model Hypothesis U as a formal statement (Tier A , DRAFT)	11
6.2	Closing verdict A5: the nonlinearity is now concrete (Tier A)	11
6.3	The conditional Millennium reduction (Tier C , DEMOTED, DRAFT)	12
7	The Four Analytical Tracks: Plan and Blocking Status	13
7.1	Current status: correctly blocked	15
7.2	Executable instrumentation delivered (Tier B)	16
8	Numerical Investigation and Its Two Failures	16
8.1	Resolution and the current evidence base (Tier C steering data)	16
9	The 3-D Bridge: Instantiating the True Nonlinearity (OP-6)	18
9.1	From computational check to kernel proof (Tier A)	19
9.2	Repairing the reduction: giving the hypotheses content (Tier C)	20
9.3	Eliminating a duplication hazard at the root	21
10	What Has Failed, and What That Taught	21
11	Summary of Verified State	22
11.1	What is genuinely established	23
11.2	What is emphatically not established	23
	Note on reference verification	24
	References	24

1 Context Management: Reading This Document and This Repository

Purpose of this section

This programme is executed collaboratively by a human mathematician and AI agents across many sessions. Context is therefore a *managed artifact*, not an accident of memory. This section specifies where authoritative context lives, how it is kept honest, and what a reader (human or machine) must load before acting. It is placed first deliberately: acting on this programme without this context has already caused documented failures (§10).

1.1 The governing documents and their precedence

Context is deliberately *not* held in any individual's or agent's memory. It is externalised into five documents with a strict precedence order:

Document	Role	Authority
SPEC.md	Normative rulebook	Highest. Defines the tier system, the honesty clause, the obstruction ledger. Where any other document conflicts with it, SPEC.md wins.
PLAN.md	Operational law	Task list with Definition-of-Done criteria and the escalation protocol (E-1...E-5). Wins over calendar/roadmap.
LEDGER.md	Claim inventory	Normative: a claim not listed here has no tier and may not be cited.
LL.md	Lessons learned	Why each process rule exists, with the incident that produced it.
CLAUDE.md	Build/verify manual	Commands, architecture, Mathlib location, working conventions.

1.2 The epistemic charter

The organising principle, stated in SPEC.md §0, is a division of labour:

The machine verifier checks the proofs; the human mathematician audits the questions.

Machine verification establishes that a proof is valid. It cannot establish that a *statement* means what the mathematics intends — that a formalisation is *adequate* to its informal gloss. Adequacy is reserved to human audit. Consequently **no LLM output gates a tier promotion**, and several artifacts in this repository are compiled, kernel-clean, and nonetheless explicitly marked DRAFT pending human statement-adequacy audit (§6).

1.3 The three-tier gating system

Tier	Meaning	Mechanically enforced gate
Tier C	Conjecture, analogy, physical narrative, unverified reduction	None, but must be tagged. Confined to <code>exploration/</code> or <code>docs/narrative/</code> . Floating-point code lives here and <i>only</i> here.
Tier B	Checkable: identities validated in exact rational arithmetic	Harness exits 0. <code>fractions.Fraction</code> or <code>int</code> only — floats are banned. Every harness ships a negative control demonstrated to actually fail.
Tier A	Established: Lean 4 kernel-compiled	Zero <code>sorry</code> ; axiom footprint naming <i>no axiom outside</i> <code>{propext, Classical.choice, Quot.sound}</code> . A strict subset is cleaner and passes — see §10, where the earlier literal-equality reading is recorded as a defect.

Two subtleties are load-bearing and were learned the hard way:

1. **“No sorry in the source” is not the Tier A gate.** A failed proof still defines its theorem name in the environment; only `#print axioms` reveals a stray `sorryAx`. The gate is the *axiom footprint*, checked on the compiled artifact.
2. **A checker that cannot fail is not a checker.** Every Tier B harness must contain a negative control that is *demonstrated* to fail when the perturbation is applied. This report documents cases where an intended negative control turned out to be undetectable by construction, and had to be redesigned (§9).

1.4 The honesty clause and the obstruction ledger

Every public claim carries its tier. Tier A status attaches to the *formal statement actually proven*, never to its informal gloss. Additionally, `SPEC.md` §1.3 requires every proposed strategy — human or agent-generated — to file an *Obstruction Compliance Note* addressing five known barriers (Table 1). Strategies that cannot are rejected at review, before any formal work begins.

ID	Obstruction	Compliance question the strategy must answer
O1	Tao’s averaged NSE blow-up	Which structural feature of the <i>true</i> nonlinearity, beyond the energy identity, does the strategy use?
O2	Supercriticality of enstrophy	Where exactly does the strategy gain over the naive enstrophy inequality, <i>uniformly</i> in α' ?
O3	Caffarelli–Kohn–Nirenberg	If the strategy implies something stronger than CKN, what new mechanism pays for it?
O4	Buckmaster–Vicol non-uniqueness	In which class is the $\alpha' \rightarrow 0$ limit taken, and why is the limit the physical solution?
O5	The Euler test	Point to the step that fails at $\nu = 0$. A strategy that would also “prove” 3D Euler regularity is presumptively wrong.

Table 1: The obstruction ledger (`SPEC.md` §1.3). These are the barriers any credible attack must address; they are stated here because they constrain everything in §7. **Re-scoped 2026-08-13** by audit verdict D2, which found that merely *listing* these barriers against a one-dimensional sequence model overstates contact with them: O1 (Tao) and O5 (the Euler test) remain meaningful for the dyadic target; O2–O4 are inherited only if and when the 3-D bridge is built.

1.5 Reproducing the verification state

The only thing that counts as “it works” is the verification pipeline:

```
./scripts/verify.sh  # Gate 1  : Tier B exact-arithmetic harnesses
                    # Gate 1b : LEDGER consistency + single-active-file lint
                    # Gate 2  : Lean kernel compile + axiom-footprint membership test
```

Every dataset in `data/` ships a `.meta` sidecar recording the exact generating command, tool versions, and `sha256sum`, so results can be audited rather than merely re-read.

Gate 1b (added 2026-08-13). `LEDGER.md` is normative — a claim not listed there may not be cited — so an *orphan* row, one citing a theorem or file that no longer exists, silently degrades the whole inventory. Gate 1b checks that every repo path cited in the ledger exists, that every Lean name in a Tier A row is actually declared, and that no banned versioned duplicates (`_v2`, `_final`) have appeared. It carries its own negative controls, run on every invocation: a dangling path, a dangling name, and a banned filename must each be *rejected*, and a clean case must be *accepted*. It found one real defect on its first run (a row citing the non-identifier fragment `_half` rather than `step1_flux_bound_half`).

A latent reproducibility defect, found by doing the cold build

`lean_src/` was made a standalone Lake project on 2026-08-13, so Gate 2 no longer depends on an external Mathlib checkout this repository does not control. (That coupling was not hypothetical: an unrelated build running in that shared checkout rebuilt its `.olean` files mid-run and failed Gate 2 on a file this repository had not touched.)

Doing the build surfaced something more serious than the inconvenience that motivated it. The project’s `lakefile` declared its dependency as `require mathlib from git with no revision pinned`. A cold build from that specification resolves to Mathlib *master* — which in general does not build against the pinned toolchain, and, worse, would have silently verified every result against a *different Mathlib* than the one all existing `LEDGER.md` claims were checked with. For a programme whose entire epistemic claim rests on machine-checkability, an unpinned verifier is a defect in the foundation, not in the build script. Mathlib is now pinned by revision, and the lockfile pinning every transitive dependency is tracked rather than ignored.

2 The Problem, Stated Honestly

2.1 Classical setting

The incompressible Navier–Stokes equations on the torus $\mathbb{T}^3$ are

$$\partial_t u + (u \cdot \nabla)u = -\nabla p + \nu \Delta u, \quad \nabla \cdot u = 0, \quad u(0) = u_0. \quad (1)$$

Leray established global existence of weak solutions [1]. The Clay Millennium Prize problem, as formalised by Fefferman [2], asks for a proof of *exactly one* of four statements — (A) existence and smoothness on $\mathbb{R}^3$; (B) existence and smoothness on $\mathbb{R}^3/\mathbb{Z}^3$ (the periodic case, which is the setting of this programme); (C) breakdown on $\mathbb{R}^3$; (D) breakdown on $\mathbb{R}^3/\mathbb{Z}^3$. Note that a *breakdown* result resolves the problem just as a regularity result does; this is consistent with the present programme’s stance that a negative verdict is a successful outcome. The obstruction is *supercriticality*: the natural energy bound is too weak to control the nonlinearity at small scales.

2.2 The regularised family and Hypothesis U

For $\alpha' > 0$ let $J_{\sqrt{\alpha'}}$ denote the projection onto Fourier modes $|k| \leq 1/\sqrt{\alpha'}$, and let $u^{(\alpha')}$ solve the truncated system

$$\partial_t u + ((J_{\sqrt{\alpha'}} u) \cdot \nabla) u = -\nabla p + \nu \Delta u, \quad \nabla \cdot u = 0, \quad u(0) = J_{\sqrt{\alpha'}} u_0, \quad (2)$$

which is globally smooth for each fixed $\alpha' > 0$.

Hypothesis 1 (Hypothesis U). For data u_0 and horizon T ,

$$\sup_{\alpha' > 0} \sup_{0 \leq t \leq T} \|\nabla u^{(\alpha')}(t)\|_{L^2(\mathbb{T}^3)} < \infty. \quad (3)$$

Proposition 2.1 (Millennium Reformulation; **Tier C** at paper level). *If Hypothesis U holds for all smooth divergence-free u_0 and all T , then 3D Navier–Stokes is globally regular. The chain is: energy identity $\Rightarrow L_t^\infty L_x^2 \cap L_t^2 H_x^1$ bounds $\Rightarrow$ Aubin–Lions compactness $\Rightarrow$ a strong- L^2 subsequential limit which is Leray–Hopf $\Rightarrow$ under U the limit lies in $L_t^\infty H_x^1 \hookrightarrow L_t^\infty L_x^6 \Rightarrow$ the Serrin class $2/s + 3/q = 1/2 \leq 1 \Rightarrow$ smooth [3].*

Retraction — please read before any other claim in this report

An earlier version of this report carried, in this box, the claim that “Hypothesis U is a restatement of the open core of the Millennium problem, not a reduction of its difficulty”. That claim was submitted to external expert audit on 13 August 2026 and was killed (verdict D1). It is retracted here in full.

The retraction is not that the claim was too strong about *difficulty* — it is that the claim was false about *what had been formalised*. The audit found that the programme’s Lean statements are indexed by a one-dimensional shell index $n \in \mathbb{N}$ rather than a frequency $k \in \mathbb{Z}^3$ (verdict A2: the index flattening destroys the lattice geometry, the density of states, and the triad constraint $k_1 + k_2 = k_3$); that the enstrophy weight was the exponential sequence 4^n rather than the Laplacian symbol $|k|^2$ (A3); and that the mode-interaction term B was an *unconstrained* parameter, so that certifying a bound for it was, in the auditor’s words, “authenticating a mathematically empty envelope” (A5). Taken together, what the formalisation describes is a **dyadic shell hierarchy**, not the unreduced 3-D equations. The prose of (1)–(3) above was never wrong; the claim that the *formal artifacts* restated it was.

The new target. The owner accepted the verdict in full (docs/Memo 1.md) and re-targeted the programme rather than defending the framing. The formal target is now **global regularity bounds for the truncated viscous Katz–Pavlović dyadic shell model** — not a proxy for the 3-D problem, but an object of study in its own right, and a respected active area of mathematical fluid mechanics (blow-up: Katz–Pavlović [7], Cheskidov [8]; regularity under strong dissipation: Barbato–Morandin–Romito [9]). A machine-verified regularity-or-blow-up theorem for that model would, to the authors’ knowledge, be a first. The 3-D bridge — the $\mathbb{Z}^3$ reindexing, OP-2 and OP-6 (§9) — is explicit, open future work, and is never again to be an implicit claim of the formalisation.

Why this is recorded as a success of the method and not merely as a setback. This programme’s Definition of Done (§11) does not include proving Hypothesis U. It includes obtaining human verdicts *whatever they are*. The tier system exists precisely so that a framing claim can be located, isolated, and killed without taking any kernel-verified theorem down with it — and that is what happened: every Lean file still compiles, and the retraction cost exactly the seven theorems that deserved to be demoted (§6). The full verdict table (A1–A5, B1–B5, C1, D1–D2) with dispositions is in LEDGER.md; the retraction text is normative in SPEC.md §1.2.

What survives the retraction, stated precisely

Hypothesis U (3) as an informal statement about the continuum truncated system, and Proposition 5.1 as paper-level mathematics, are unaffected — they were always **Tier C** at paper level and are so labelled above. What is retracted is the assertion that the *Lean formalisation* of §6 is a formalisation of *that* statement. Audit verdicts A1 (the quantifier order $\exists C \forall N$) and A4 (quantifying over every solution) were upheld, so the statement’s *logical shape* was sound; it was its *referent* that was wrong. Under the pivot, A2 and A3 dissolve rather than being repaired: for the dyadic shell model the index *is* the object and the weight $4^n = k_n^2$ *is* the correct weight. A5 was repaired by instantiating B (§6).

3 Dual-Scale Geometry and T-Duality

3.1 Physical motivation (**Tier C**) and its strict boundary

The programme’s geometric inspiration is T-duality in string theory. For a dimension compactified on a circle of radius R , the free-boson partition function is invariant under $R/\sqrt{\alpha'} \rightarrow \sqrt{\alpha'}/R$ with simultaneous exchange of momentum and winding quantum numbers — this is a genuine theorem about the CFT spectrum, established as Eq. (2.2.11) of [14]. The fixed point of that map is $R = \sqrt{\alpha'}$.

Boundary of the analogy — Tier C, quarantined

Reference [14] was retrieved and read for this report, and the distinction it forces is worth stating precisely, because the programme’s earlier internal documents blurred it. The **rigorous** content is the exact $R \leftrightarrow \alpha'/R$ symmetry of the spectrum. The frequently quoted claim that T-duality entails a “**minimal length**” is a *physical interpretation* of that symmetry — the argument being that a string cannot distinguish R from α'/R , so sub- $\sqrt{\alpha'}$ geometry is not resolvable — and is **not itself a separately proved theorem** in the sources consulted.

This report therefore uses T-duality **only** as motivation for the *shape* of a cutoff function. The mechanism acting on the PDE is the Fourier cutoff $J_{\sqrt{\alpha'}}$ and nothing more. Deriving the truncated dynamics from a genuine metric deformation is an *open problem* of this programme (OP-1), not an assumption. All cosmological interpretation — bounce scenarios, dark-sector mass laws, K3/mock-modular numerology — is quarantined in `docs/narrative/`, may never be imported by `lean_src/`, and may never be cited in a tier promotion.

3.2 The effective radius: what is actually proven (**Tier A**)

Define, for $\alpha, R > 0$,

$$R_{\text{eff}}(\alpha, R) := \max\left(R, \frac{\alpha}{R}\right). \quad (4)$$

The following are kernel-verified in `lean_src/CallensDualScale.lean`, with axiom footprint exactly `[propext, Classical.choice, Quot.sound]`:

Formal name	Statement	Reading
<code>Reff_pos</code>	$0 < R_{\text{eff}}(\alpha, R)$	positivity
<code>Reff_ge_sqrt</code>	$\sqrt{\alpha} \leq R_{\text{eff}}(\alpha, R)$	universal minimum scale
<code>Reff_bounce</code>	$R < \sqrt{\alpha} \Rightarrow R_{\text{eff}} = \alpha/R$	sub-cutoff scales reflect
<code>Reff_inertial</code>	$\sqrt{\alpha} \leq R \Rightarrow R_{\text{eff}} = R$	macroscopic scales unmodified
<code>Reff_tdual</code>	$R_{\text{eff}}(\alpha, \alpha/R) = R_{\text{eff}}(\alpha, R)$	self-dual symmetry
<code>Reff_gt_sqrt_of_ne</code>	$R \neq \sqrt{\alpha} \Rightarrow R_{\text{eff}} > \sqrt{\alpha}$	strictness
<code>Reff_eq_sqrt_iff</code>	$R_{\text{eff}} = \sqrt{\alpha} \iff R = \sqrt{\alpha}$	sharpness: the minimum is attained at exactly one point

Scope — what these theorems are and are not

These are theorems about the real-valued function $\max(R, \alpha/R)$. They are elementary, they are correct, and they are *not yet about fluids*. Their role is to fix the geometric language precisely; the parameter α' is carried as a *hypothesis parameter* ($0 < \alpha$), never as an axiom, because Hypothesis U itself quantifies over it. A prior version of this repository asserted these as proven when no such proofs existed; that defect and its correction are recorded in `LEDGER.md` under “Retired / corrected claims”.

3.3 The Symmetric-Square Lock (Tier A)

The programme’s second structural ingredient is an algebraic rigidity constraint. If a sequence satisfies a second-order recurrence with operator L_2 ,

$$u_{n+2} = a u_{n+1} + b u_n,$$

then $v_n := u_n^2$ satisfies the third-order recurrence

$$v_{n+3} = (a^2 + b) v_{n+2} + b(a^2 + b) v_{n+1} - b^3 v_n, \quad (5)$$

which is the symmetric square $L_3 = \text{Sym}^2(L_2)$. If L_2 has characteristic roots $\{\lambda, \mu\}$, then L_3 has roots $\{\lambda^2, \lambda\mu, \mu^2\}$, and the coefficients of (5) are precisely their elementary symmetric functions:

$$e_1 = \lambda^2 + \lambda\mu + \mu^2 = a^2 + b, \quad e_2 = -b(a^2 + b), \quad e_3 = (\lambda\mu)^3 = -b^3.$$

This spectral content is verified as `sym2_symmetric_functions` — it is what makes the identity *structural* rather than a coincidence. The symmetric square of a second-order operator being third-order is classical in the theory of linear differential operators and Picard–Fuchs equations [15].

Intended role (Tier C). The *conjectural* purpose of the lock is to tie macroscopic transport (L_3) rigidly to a microscopic fiber operator (L_2), so that high-frequency modes cannot evolve as independent degrees of freedom — the hoped-for mechanism against Tao-type “choreographed” blow-up [4]. **Its relevance to Navier–Stokes is a conjecture of this programme and is not established.** It must be earned in the dyadic laboratory (§5).

4 The “Holographic Algebra” Organising Concept

Tier and status

The name *HoloAlg* (Holographic Algebra) denotes the programme’s organising picture, not a proven theory. **Everything in this section is Tier C** except where it explicitly cites a Tier A result from §3. It is included because a reader must understand the motivating picture to follow why the four tracks (§7) were chosen — not because the picture carries evidential weight.

4.1 The picture

The organising idea combines two elements already made precise:

1. **A dual-scale geometry** in which the effective scale seen by the fluid is $R_{\text{eff}}(\alpha', R) = \max(R, \alpha'/R)$, so that (i) macroscopic scales are unmodified (`Reff_inertial`: “inertial invisibility”), and (ii) no scale below $\sqrt{\alpha'}$ is ever resolved (`Reff_ge_sqrt`). The regularisation is thereby confined strictly to the sub-cutoff regime.
2. **An algebraic lock** (§3) constraining macroscopic modes to be generated by products of microscopic ones, restricting the admissible spectrum to $\{\lambda^2, \lambda\mu, \mu^2\}$.

The label “holographic” is used in the loose sense that macroscopic (bulk) transport is determined by — and rigidly constrained by — data on a lower-dimensional algebraic structure (the fiber operator’s spectrum). **No claim is made of a relationship to the AdS/CFT correspondence or any holographic duality in the technical sense of theoretical physics.** No such correspondence has been constructed here, and none is used.

4.2 Why this picture would matter, *if* it could be substantiated

For fluid mechanics, the mechanism would matter because it addresses obstruction O1 (Table 1) with a *structural* feature rather than an energetic one: Tao’s averaged system satisfies the energy identity yet blows up, so any successful strategy must use something beyond energy. Algebraic rigidity of the mode spectrum is such a candidate. For astrophysics and cosmology, the same cutoff geometry has been *speculatively* associated with bounce scenarios — but that association is quarantined narrative, carries no evidential weight here, and is deliberately excluded from the mathematical core.

The programme’s own structural finding against itself

When the four tracks’ missing definitions were drafted (`docs/designs/TRACK_DEFINITIONS_DRAFT.md`), the drafting process surfaced a negative structural result about the programme’s own architecture: under the natural choice of entropy functional, **Track T3 collapses into the direct production-identity attack rather than being independent**, and **T1 and T4 both depend on the same unproven embedding (OP-2)**, so agreement between them would be correlated rather than independent evidence. The advertised “four independent attacks” is therefore *not* currently four independent attacks. This finding is recorded rather than suppressed.

5 The Dyadic Shell Model: Established Results

This section was written when the truncated dyadic shell model was the programme’s *laboratory* — a place to earn techniques before touching the continuum PDE. After the audit of 13 August

2026 it is the programme's *target* (§2). The results below are unchanged by that reclassification; what changed is that they are no longer preliminaries to something else. The model is due to Katz and Pavlović [7], whose published form (verified by retrieval, as restated in [8]) is $\frac{d}{dt}u_n + \nu 2^{2\alpha n}u_n - \lambda^n u_{n-1}^2 + \lambda^{n+1}u_n u_{n+1} = g_n$. The system used here is the unforced case with λ^n playing the role of k_n : shells $n = 0, \dots, N$ carry wavenumbers $k_n = 2^n$ and real amplitudes a_n , with

$$\frac{da_n}{dt} = \text{NL}_n - \nu k_n^2 a_n, \quad \text{NL}_n := k_{n-1} a_{n-1}^2 - k_n a_n a_{n+1}, \quad (6)$$

with truncation conventions $a_{-1} = 0$, $a_{N+1} = 0$. Energy is $E = \frac{1}{2} \sum a_n^2$ and enstrophy $\Omega = \frac{1}{2} \sum k_n^2 a_n^2$.

Scope note (honesty clause)

This model is **not** the Navier–Stokes equations. Its nonlinearity is **never** asserted in this repository to be derived from, or to approximate, $(u \cdot \nabla)u$; the audit of 13 August 2026 (§2) exists precisely because an earlier framing let that boundary blur. Results below are exact statements about finite sums and their consequences; no existence theory for the ODE system is invoked. What the model *is* is a well-studied object with a real literature and real open questions — which is what makes it a legitimate target rather than a proxy.

5.1 Energy flux telescoping and monotonicity (Tier A)

In `lean_src/DyadicShells.lean`: the nonlinear energy flux telescopes for *arbitrary* k and a (`dyadic_flux_telescopes`); it vanishes under the truncation $a_{N+1} = 0$ (`dyadic_flux_zero_of_boundary`); and consequently $dE/dt \leq 0$ for $\nu \geq 0$ (`energyRate_nonpos`). Non-vacuity is witnessed: a concrete instance pins $dE/dt = -1$ exactly, so the inequality is not the trivial $0 \leq 0$.

5.2 The exact enstrophy-production identity (Tier A)

The programme's most substantive dyadic result concerns *enstrophy*, the quantity Hypothesis U is about. Unlike energy flux, the enstrophy production sum does *not* telescope for arbitrary k ; it collapses only under the doubling hypothesis $k_{n+1} = 2k_n$, and it collapses not to zero but to a multiple:

Theorem 5.1 (Dyadic enstrophy production; `enstrophy_production_dyadic`). *Under $k_{n+1} = 2k_n$ and truncation $a_{N+1} = 0$,*

$$\sum_{n=0}^N k_n^2 a_n \text{NL}_n = 3 \sum_{n < N} k_n^3 a_n^2 a_{n+1}, \quad (7)$$

so that the enstrophy evolves as

$$\frac{d\Omega}{dt} = 3 \sum_n k_n^3 a_n^2 a_{n+1} - \nu \sum_n k_n^4 a_n^2. \quad (8)$$

This is the dyadic analogue of vortex stretching, made exact: a signed production term competing against dissipation. The coefficient 3 is not fitted — it is $r^2 - 1$ for doubling ratio $r = 2$, and the Tier B harness independently confirms the general formula at $r = 3$ (coefficient 8). A kernel-verified negative control (`negative_control_nondoubling`) establishes that the identity genuinely *fails* for non-doubling k ; the checker can fail.

5.3 A local production bound (Tier A)

Writing $S_N := \sum_{n < N} k_n^3 a_n^2 a_{n+1}$ and $\Omega_N := \frac{1}{2} \sum_{n \leq N} k_n^2 a_n^2$, `lean_src/EnstrophyProductionBound.lean` proves, in a sqrt-free three-step chain (elementary algebra, $\sum x^2 \leq (\sum x)^2$ for nonnegative x , Cauchy–Schwarz):

$$S_N^2 \leq 2\Omega_N^3. \quad (9)$$

Precise scope of (9) — this is where over-reading would be easy

This is a *local, instantaneous, algebraic* bound. It does *not* use dissipation, it does *not* integrate in time, and it does *not* address uniformity in N . Note the exponent: production is bounded by $\Omega^{3/2}$, which is exactly the supercritical scaling that makes the enstrophy differential inequality $\dot{\Omega} \lesssim \Omega^{3/2}$ insufficient for global control. **Inequality (9) is therefore consistent with blow-up and is not evidence for Hypothesis U.** It is infrastructure: an exact, reusable, verified inequality, nothing more. The Tier B mirror further found, honestly, that the negative control breaks Step 1 under non-doubling k but does *not* break the final bound on the tested state family — an unforced negative result, recorded.

6 Statement-Level Formalisation, the Pivot, and the Demotion

6.1 The shell-model Hypothesis U as a formal statement (Tier A, draft)

`lean_src/DyadicShell_Statements.lean` (renamed 2026-08-13 from `HypothesisU_Statements.lean`, Memo Task 1) gives a non-vacuous Galerkin-shape formalisation. Writing $u : \mathbb{R} \rightarrow (\mathbb{N} \rightarrow \mathbb{R})$ for mode amplitudes, a solution at cutoff N satisfies: (i) $u(0) = \text{truncate } N u_0$; (ii) for every t and every retained mode $n \leq N$, $\frac{d}{ds} u(s, n)|_{s=t} = B_n(u(t)) - \nu w_n u(t, n)$; (iii) $u(t, n) = 0$ for $n > N$. Then

$$\text{HypothesisU} \equiv \exists C > 0. \forall N. \forall u. [u \text{ solves at cutoff } N] \Rightarrow \forall t \in [0, T]. \Omega_N(u, t) \leq C. \quad (10)$$

Why the quantifier order is the entire content

$\exists C \forall N$ — the constant is chosen *before* the cutoff. Swapping to $\forall N \exists C$ would be the trivial statement that each finite-dimensional truncation is bounded on its own, which carries no information about the $\alpha' \rightarrow 0$ limit. That swap is the failure mode this ordering exists to exclude.

Two prior defects were repaired by explicit human decision (recorded as Q1, Q2): initial data must be *projected* per cutoff, else the hypothesis set is empty at low N and the statement holds vacuously; and the weight is fixed concretely to $w_n = (2^n)^2$, so that “enstrophy” is provably a nonnegative sum of squares rather than merely a name. The file also *machine-checks the refutation* of the prior formalisation: quantifying over all smooth fields, without the equation as a constraint, yields a statement that is **provably false** (`unconstrained_bound_false`), since scaling a field by c scales enstrophy by c^2 .

6.2 Closing verdict A5: the nonlinearity is now concrete (Tier A)

Until 13 August 2026 the mode-interaction term B in (10) was an **abstract parameter**, and the previous version of this report named that the decisive open item. The audit put it far more sharply (verdict A5): certifying a bound for an unconstrained B is a “fatal flaw . . . authenticating a mathematically empty envelope” — an arbitrary truncated sequence system simply blows up, so the envelope contains nothing. Under the pivot this was repaired directly, by instantiating B as the Katz–Pavlović shell nonlinearity itself (Memo Task 2):

$$\text{shellB} : B_0(v) = -(k_0 v_0 v_1), \quad B_{m+1}(v) = k_m v_m^2 - k_{m+1} v_{m+1} v_{m+2}. \quad (11)$$

The bottom shell has no influx, so k_{-1} is never evaluated; for $n \geq 1$ this is exactly NL_n of (6). The structural property whose absence the audit called fatal is now a kernel-checked theorem, obtained from an exact telescoping identity:

Theorem 6.1 (Shell energy pairing; `sum_mul_shellB`, `shellB_energy_conservation`). *For all $k, v : \mathbb{N} \rightarrow \mathbb{R}$ and all N , writing $\text{out}_m := k_m v_m^2 v_{m+1}$, each summand satisfies $v_n B_n(v) = \text{out}_{n-1} - \text{out}_n$ (reading $\text{out}_{-1} = 0$), so that*

$$\sum_{n \leq N} v_n B_n(v) = -k_N v_N^2 v_{N+1}, \quad \text{and hence} \quad v_{N+1} = 0 \implies \sum_{n \leq N} v_n B_n(v) = 0.$$

The truncation hypothesis $v_{N+1} = 0$ is clause (iii) of `IsGalerkinSolution`, so this transfers to solutions with no extra assumption: `galerkin_shellB_conservation` states that every Galerkin solution of the concrete model conserves energy exactly, at every time. The viscous term of `IsGalerkinSolution` is $-\nu w_n u_n = -\nu k_n^2 u_n$, which is already the shell model’s dissipation verbatim; together with (11) this makes

$$\text{DyadicShellHypothesisU} := \text{HypothesisU} \nu T (\text{shellB } k) (n \mapsto k_n^2) u_0$$

the programme’s headline statement, in which **nothing abstract remains**. It is a statement about precisely the truncated viscous Katz–Pavlović system — the same model whose exact enstrophy-production identity is proven in §5.

What this does and does not fix

It fixes A5, and only A5. `DyadicShellHypothesisU` is a *statement*; it is not proven, and nothing here is evidence that it is true. The abstract- B scaffolding is retained in the file solely because the refutation theorems and the inhabitation witnesses quantify over it; it carries no headline claim. And instantiating B with the *true 3-D* nonlinearity remains untouched and open — that is OP-6 (§9), now correctly labelled as the 3-D bridge rather than as this file’s missing piece.

6.3 The conditional Millennium reduction (**Tier C**, demoted, draft)

`lean_src/MillenniumReduction.lean` encodes the logical shape of Proposition 5.1 with the two undischarged analytic steps as *named hypothesis parameters*, never axioms:

`millennium_reduction : ( $\forall T$ , HypothesisU)  $\rightarrow$  AubinLionsStatement  $\rightarrow$  ProdiSerrinStatement  $\rightarrow$  GlobalRegu`

The proof is the one-line composition `hPS (hAL hU)`. The intent was that all mathematical weight be parked in the hypotheses’ *types*, so the Lean kernel would track the outstanding debt in the type signature rather than in the axiom environment — the pattern prescribed after a prior version of this work carried a raw `axiom aubin_lions_compactness`.

An adversarial self-review against `SPEC.md`’s actual text found three statement-adequacy gaps, all repaired by explicit decision (Q3–Q5): global regularity had covered only *time*-regularity, and now additionally requires spatial smoothness via membership in every Fourier-side Sobolev class H^s ; the theorem had fixed a single horizon T while concluding a T -independent result, and now quantifies over all T ; and the weight nonnegativity $w \geq 0$ is now an explicit hypothesis.

Demotion to — audit verdict B1

That intent did not survive audit. `AubinLionsStatement` and `ProdiSerrinStatement` are bare `Prop` parameters, and the auditor’s finding (B1) is that bare `Prop → Prop` arrows “completely bypass PDE theory. The Lean kernel is merely verifying a logical tautology ($A \rightarrow B \rightarrow C$)” — which violates the Definition of Done’s requirement of bona fide statements. **The verdict is accepted and the demotion is executed: all 7 theorems in this file are Tier C as of 2026-08-13.** They remain kernel-compiled and wired into Gate 2 so that they cannot silently rot, but they may no longer be cited as Tier A.

The distinction being drawn is sharp and worth stating plainly: parking analytic debt in a hypothesis’s *type* is only honest if that type has content. Here it had none — a `Prop` variable constrains nothing — so the type signature recorded the debt in name only. The repair is Memo Task 4: import or define the required sequence spaces (ℓ^2 , ℓ^p), state the Aubin–Lions and Prodi–Serrin analogues *for the shell model* with their actual topological content, and (audit item B4) specialise the reduction to the concrete object rather than to a generic B . That is the next design task and is **open**; `PLAN.md` §10 records that no implementation may begin before the design memo exists.

One audit item was repairable immediately (B2, revise). A previous independent review had raised, and this report previously recorded as an open candidate issue, whether the untruncated solution may differ per horizon T without a coherence clause. The audit confirmed it: the existential had to be hoisted. It now is — `HasGlobalBoundedLimit` fixes *one* limit trajectory u_{lim} across *all* horizons T , rather than permitting a different limit for each. The Cantor diagonal argument that would actually produce such a trajectory is thereby relocated into `AubinLionsStatement`’s type, where the audit said it belongs, instead of being silently absent; formalising the diagonalisation itself is part of the Task 4 repair.

Audit status of the two files after the pivot

`DyadicShell_Statements.lean` is **Tier A** and DRAFT: it compiles clean with correct axiom footprints, its quantifier shape was upheld (A1, A4), its nonlinearity is now concrete (A5), but it has **not** passed a full human statement-adequacy audit. `MillenniumReduction.lean` is **Tier C** and DRAFT pending the Task 4 repair. The general lesson is the one this section now demonstrates twice over: fixing the gaps found in a self-review (Q3–Q5) did *not* certify that no others remained, and the ones that remained were the structural ones. Authorship — including adversarial self-review — never licenses a claim. Only human audit does, and when it was finally obtained it removed a claim rather than confirming one.

7 The Four Analytical Tracks: Plan and Blocking Status

Stage 3 of the programme proposes to prove Hypothesis U *scale-by-scale* rather than as a monolith, partitioning Fourier space into dyadic shells and controlling $\sum_j E_j^{(\alpha')}(t)$. Four deep analytical disciplines are proposed as independent attacks. Each is an *analogy with a falsification plan*, and each carries a kill criterion agreed in advance.

Status of this section after the pivot

All four tracks are formulated against the *continuum* problem on $\mathbb{T}^3$ and depend on $\mathbb{Z}^3$ lattice structure (T1’s triads, T4’s percolation field). They therefore sit **downstream of the 3-D bridge** (§9), not of the programme’s current formal target. They are retained here unchanged, and unchanged also in their status: blocked. Their blockage was never contingent on the framing that was retracted — it is contingent on four objects that do not exist, listed below.

	Discipline	Proposed mechanism	First falsifiable milestone / kill criterion
T1	Bourgain–Demeter ℓ^2 decoupling [10]	Arithmetic depletion: the Sym ² constraint starves triadic resonances $k_1 + k_2 = k_3$	Exact count of constrained vs. unconstrained resonant triads. Kill: constrained count grows at the same order $\Rightarrow$ no depletion.
T2	Mouhot–Villani non-linear Landau damping [11]	Phase mixing suppresses “enstrophy echoes” via a Newton-scheme argument in a high-regularity class	A precise <i>definition</i> of an enstrophy echo in the truncated system, plus its exact evolution. Kill: echoes not suppressed even in the dyadic model.
T3	Golse–Saint-Raymond hydrodynamic limits [12]	A relative-entropy functional controls dissipation uniformly in α'	A relative-entropy inequality for the truncated system at fixed α' . Kill: inequality already fails at fixed α' .
T4	Duminil-Copin percolation [13]	High-enstrophy regions form <i>subcritical</i> percolation clusters $\Rightarrow$ singular set of dimension 0	Recover the CKN dimension bound [5] in the lattice formulation. Kill: cannot recover even CKN.

A near-miss worth recording: the correction that would itself have been an error

An earlier draft of this report flagged an apparent defect in the programme’s specification. The spec describes T2 as resting on *Gevrey-class* regularity; a first retrieval pass indicated that Mouhot–Villani’s *core* result is set in *analytic* regularity, and the obvious move was to “fix” the spec by substituting one for the other.

That fix would have been wrong. Re-fetching the abstract directly — rather than relying on the intermediate summary the first pass produced, parts of which were flagged as paraphrase — gives *both*: the paper establishes “exponential Landau damping in **analytic** regularity”, while the version note records that the main result “now covers Coulomb and Newton potentials, and (2) some classes of **Gevrey** data”. The base theorem is analytic; the final version does cover Gevrey. **The spec was defensible as written and was left standing.**

The methodological point is the reason this is in the report at all: the first pass had enough material to justify editing the normative rulebook, and editing it would have *introduced* an error on the strength of a paraphrase — exactly the failure mode rule LL-6 exists to prevent, one level up from where it was first learned. **When a proposed correction rests on a summary rather than a primary quotation, flag it and re-fetch; do not edit the rulebook.**

One substantive caveat did survive, and it is the part that bears on feasibility: coverage is “*some* classes of Gevrey data”, not all. Any T2 argument leaning on Gevrey regularity must identify *which* class and must not assume the full Gevrey scale.

A disanalogy that has not been resolved away

The Golse–Saint-Raymond programme [12] concerns the Boltzmann-to-Navier–Stokes hydrodynamic limit, whereas this programme’s $\alpha' \rightarrow 0$ limit is **not** a Knudsen limit and has no kinetic layer — a disanalogy flagged internally (review finding M4) and repeated here so T3 is not over-read. This one stands.

7.1 Current status: correctly blocked

All four tracks are blocked on undefined objects — and this is the correct state

Tracks T1–T4 require four objects that **do not exist** in this programme or, to the authors’ knowledge, in the literature:

- **OP-2**: the Sym^2 -constrained spectrum on T^3 . The lock is proven for *scalar recurrences*; its embedding into 3D Fourier dynamics indexed by $k \in \mathbb{Z}^3$ is open. Any invented embedding would make T1’s count meaningless.
- **OP-3**: “enstrophy echo” as a formula. Landau-damping echoes are defined in Vlasov phase space; no NSE analogue exists.
- **OP-4**: the entropy functional h and reference state $\bar{u}$. Choosing h is the mathematical content of T3.
- **OP-5**: the coupling of the percolation field to Sym^2 . Without it, T4 measures generic percolation, which cannot beat CKN.

Under escalation rule E-1, an agent encountering a missing definition must **stop and escalate**, never improvise. Draft definitions have been authored and await human audit; *authorship never unblocks a track — only audit does*. Progress here is deliberately gated on human judgement rather than machine throughput.

7.2 Executable instrumentation delivered (Tier B)

Definition-independent tooling has nonetheless been built and certified, so that the tracks can execute immediately once unblocked: exact lattice representation counts $r_3(n)$ for $n \leq 10^4$ with Legendre anchors; exact unconstrained triad counts for $|k_i|^2 \leq M^2$; and an exact 3D periodic site-percolation toolkit (union-find over integers, wrap detection, 27 checks, negative control that fails when periodicity is dropped).

8 Numerical Investigation and Its Two Failures

The central empirical question is whether $\sup_t \Omega_N(t)$ grows with the cutoff N or saturates. Two successive designs *failed their feasibility premises*, for genuinely different reasons; both are documented rather than quietly replaced.

1. **Explicit float RK4 (D4).** Stiffness: stability requires $\Delta t \sim 1/(\nu k_N^2)$, which collapses as N grows. 21 of 36 configurations were computationally infeasible.
2. **Exact-rational IMEX-Euler (D5).** The implicit viscous treatment removed the stiffness barrier as designed — but the scheme failed anyway, for an unrelated reason: iterating a *nonlinear* rational map in exact arithmetic causes the numerators and denominators to grow geometrically in bit-length (hundreds of digits within 2–6 steps), while the state magnitude stays small. **This is a representation blow-up, not a numerical instability.** Implicit-versus-explicit was the wrong axis on which to redesign.

Generalised lesson, now a standing rule

Exact rational arithmetic is excellent for closed-form algebraic identities and inequalities proven in one shot, and poor for *iterating* a nonlinear map over many steps. For trajectories, use dual-precision floating point (Tier C steering) and reserve exactness for the algebra. After two failed designs, the methodology decision was escalated to the human owner rather than a third variant being chosen unilaterally.

8.1 Resolution and the current evidence base (Tier C steering data)

The accepted resolution was to stop pursuing exact-rational certification and retain dual-precision (float64 + 50-digit `mpmath`) runs as permanent Tier C steering data. Across the full grid $N \in \{8, 12, 16, 20, 24\}$, several low-viscosity configurations initially reported divergence, with the two precisions agreeing to $\sim 10^{-14}$ — ruling out rounding, but *not* distinguishing genuine divergence from an under-resolved step size (the design's Δt formula never depended on ν). A step-size refinement study resolved this:

Finding	Detail
Divergences are artifacts	All four diverging (ν , profile) pairs flip to OK at a finite refinement level and stay OK at every finer level, with $\sup \Omega$ <i>decreasing</i> as $\Delta t \rightarrow 0$ (e.g. $45.1 \rightarrow 21.8 \rightarrow 17.3 \rightarrow 15.8$) — the signature of a discretisation artifact resolving, not a fixed-time blow-up.
N -independence	44 of 45 configurations are <i>exactly</i> flat in N (bit-identical $\sup \Omega$ at every cutoff). This was reported as striking evidence. It is not evidence at all — see the retraction immediately below.
The one exception, explained	$\nu = 0.1$, profile P2 appeared to grow linearly in N (4.5, 6.5, $\dots$, 12.5). Traced by hand: this equals $\frac{1}{2}(N + 1)$, which is profile P2's <i>analytic initial</i> enstrophy — P2 excites every retained mode by construction, so a larger cutoff starts with more enstrophy independent of any dynamics. The trajectory collapses immediately from $t = 0$ and decays thereafter. An initial-data artifact, not dynamical growth.

Retraction (14 August 2026): the N -flatness is an artifact of the grid

An earlier version of this report presented the flatness above as the evidence base for a uniformity verdict. **That reading was wrong, and the error was the report's own.** Direct measurement of the shell-amplitude profile shows the cascade never reaches the shells the grid varies. At $\nu = 10^{-3}$ the enstrophy is carried almost entirely by a *single* shell:

shell	peak $ a_n $	$k_n^2 a_n^2$	share of $2\Omega \approx 140$
$n = 8$	4.5×10^{-2}	132.7	$\approx 95\%$
$n = 10$	1.7×10^{-6}	3.0×10^{-6}	2×10^{-8}

The grid's *smallest* cutoff was $N = 8$ — exactly one shell above the decisive one. Raising N from 8 to 24 appended shells contributing $\sim 10^{-8}$ relative, which is why 44 of 45 configurations agreed *bit-identically*: adding numerically zero shells changes $\sup \Omega$ by nothing. **The experiment varied the one parameter along the one range in which it could have no effect.** Lowering to $N = 7$ would have removed 95% of the enstrophy.

“Flat in N ” is therefore close to a tautology here and cannot distinguish “*the bound is uniform in the cutoff*” from “*no regime was tested in which the cutoff binds*”. No verdict may be issued on this evidence base, and none has been.

The failure mode is now mechanically excluded

A Tier B harness (`tests/tier_b_grid_adequacy.py`, wired into Gate 1) rejects such a grid *before it runs*, on exact integer arithmetic. Balancing viscous against nonlinear rates with the K41 amplitude $a_n \sim k_n^{-1/3}$ gives $k_d \sim \nu^{-3/4}$; with $k_n = 2^n$ the cutoff lies at or above the dissipation scale iff $2^{4N}\nu^3 \geq 1$, i.e. writing $\nu = p/q$, iff the integer inequality $2^{4N}p^3 \geq q^3$ holds. No floats, no logarithms.

Its negative control is the historical grid itself: **0 of 15 configurations had a biting cutoff.** The redesigned measurement (`docs/designs/D6_CUTOFF_BITING_MEASUREMENT.md`) sweeps N *across* the dissipation shell instead of above it, reports the production flux at the cutoff and the shell-population profile alongside $\sup \Omega$, and carries a $\nu = 0$ control where blow-up is a published theorem — the sharper control, because it is the one this programme's hoped-for conclusion would predict wrongly.

9 The 3-D Bridge: Instantiating the True Nonlinearity (OP-6)

This section describes work aimed at the *true* 3-D nonlinearity. It should now be read as work on the **bridge** between the programme’s formal target (the shell model, §5, §6) and its motivating problem (§2) — open future work, and not a claim that the bridge exists. Note the distinction from §6: B in (10) is no longer abstract, having been instantiated by the shell nonlinearity (11); what remains open is instantiating it by the Fourier–Galerkin nonlinearity on $\mathbb{Z}^3$, which is a different and strictly harder object.

A repository-wide search established that **no Fourier-triad convolution formula exists anywhere in this programme’s sanctioned content** — the triad tooling counts lattice points without amplitudes; the OP-2 draft self-admits it discards the angular structure of k “which is precisely where 3-D vortex-stretching geometry lives”; and the dyadic NL_n is never asserted to derive from $(u \cdot \nabla)u$. Instantiating B by the 3-D nonlinearity is therefore a genuine E-1 blocker.

On $\mathbb{T}^3$, with $u = \sum_k \hat{u}_k e^{ik \cdot x}$, $k \cdot \hat{u}_k = 0$, $\hat{u}_{-k} = \overline{\hat{u}_k}$, and Leray projector $P(k) = I - \frac{k \otimes k}{|k|^2}$, the truncated nonlinearity is

$$\partial_t \hat{u}_k = -\nu |k|^2 \hat{u}_k - i \sum_{\substack{p+q=k \\ p,q \in \Lambda}} P(k) [(q \cdot \hat{u}_p) \hat{u}_q]. \quad (12)$$

An erratum caught by exact arithmetic — a case study in why the method exists

The formula first recorded in the scoping memo, taken from a single literature cross-check, read $P(k)[q(\hat{u}_p \cdot \hat{u}_q)]$ — the *vector* q times the scalar velocity–velocity dot product. The Tier B harness then *proved*, in exact arithmetic, that this expression is **identically zero for every k** : $P(k)$ annihilates anything parallel to k , and since every pair (p, q) with $p + q = k$ appears alongside its swap (q, p) , the two terms cancel exactly, $P(k)[qs] + P(k)[ps] = P(k)[(p + q)s] = P(k)[ks] = 0$. That cannot be the real nonlinearity. The corrected form (12) — q dotted with $\hat{u}_p$, a scalar, times the vector $\hat{u}_q$ — matches the Fourier transform of the convective derivative and is nonzero. The error was caught by computation, not by inspection.

The resulting harness (`tests/tier_b_nse_triad_convolution.py`, **Tier B**) certifies two structural facts in exact Gaussian-rational arithmetic at $M = 1, 2, 3$: *transversality* $k \cdot N(u)_k = 0$ (unconditional, since $P(k)$ projects onto $k^\perp$), and *detailed energy conservation* $\sum_k \overline{\hat{u}_k} \cdot N(u)_k = 0$ given divergence-free, conjugate-symmetric input. Two negative controls — dropping the Leray projection, and breaking the divergence-free condition on one mode — each break exactly one fact and leave the other intact, matching the hand derivation rather than merely “changing something”.

Recorded honestly: several negative controls were undetectable by construction

Two natural “plausible bug” perturbations turned out to be invisible: swapping $p \leftrightarrow q$ throughout (the pair set is symmetric under relabelling for *every* fixed k , not merely after summing), and substituting the outer k for the summed q (the difference vanishes because $p \cdot \hat{u}_p = 0$ by divergence-freeness). The negative control actually used — multiplying by the wrong velocity factor — was found by *testing* candidates, not by assuming one would work. Additionally, the general triad identity underlying energy conservation is here verified *computationally* at three truncations, **not proven symbolically**; a Tier A proof remains future work.

9.1 From computational check to kernel proof (Tier A)

The harness above certifies two facts *computationally*, at three truncations, and its own docstring declines to claim either as proven. Both are now Tier A, in `lean_src/AbstractAlgebraicConservation.lean` — renamed 2026-08-13 from `TriadConservation.lean` under audit verdict C1, which found the old name OVERSTATES REACH: the file’s abstraction bypasses geometric aliasing on $\mathbb{Z}^3$ entirely, and the name should say so. The docstring now states that exclusion explicitly:

Theorem 9.1 (Termwise triad cancellation; `triad_pairing`). *If $k_p + k_q + k_r = 0$ componentwise and the field is divergence-free at p ($k_p \cdot u_p = 0$), then the two orderings of the triad summand cancel exactly:*

$$(k_q \cdot u_p)(u_q \cdot u_r) + (k_r \cdot u_p)(u_r \cdot u_q) = 0.$$

Why the first attempt failed, and what actually worked

The harness records that a direct *three-way* relabelling (cycling $p \rightarrow q \rightarrow r$) did not close. It does not need to. The cancellation is **two-way**, and stronger than a summation argument: it is **termwise**. Pairing (p, q, r) with (p, r, q) — swapping only the last two indices, leaving p fixed — and using symmetry of the pairing followed by additivity,

$$[(k_q \cdot u_p) + (k_r \cdot u_p)](u_q \cdot u_r) = ((k_q + k_r) \cdot u_p)(u_q \cdot u_r) = (-k_p \cdot u_p)(u_q \cdot u_r) = 0.$$

This was re-derived by hand and confirmed on 6486 triples in exact arithmetic *before* any Lean was written. Summing it over any finite index set closed under the swap gives detailed energy conservation (`triad_sum_zero`); transversality (`transversality_of_sum`) is separate and unconditional, following from the projector landing in $k^\perp$.

Scope: this does not close OP-6

The file is stated over an *arbitrary commutative ring*, with an arbitrary additive index group and additive wavevector map. It is deliberately **not** stated over $\Lambda \subset \mathbb{Z}^3$ with complex velocities, because that concrete apparatus *is* the open decision D1. **The bridge from this lemma to the concrete Fourier–Galerkin setting is not built**; nothing here instantiates B ; nothing here is a statement about Navier–Stokes solutions. One hypothesis is load-bearing rather than technical: the swap has fixed points (where $q = r$, forcing $p = -2q$) at which the pairing yields only $2f = 0$, so the ring must have no 2-torsion. For the intended instantiations $(\mathbb{R}, \mathbb{C})$ this holds.

Open decisions. Instantiating B is not a substitution but a restructuring: the state representation must move from $\mathbb{N} \rightarrow \mathbb{R}$ shell scalars to $\mathbb{Z}^3 \rightarrow \mathbb{C}^3$ Fourier vectors carrying divergence-free and reality constraints. Whether to undertake that full reindexing or seek a smaller shell-preserving proxy, and whether to sequence it before or after OP-2, are open questions reserved to the human owner.

A second, sharper limit on the pivot (14 August 2026)

The audit established that the formalisation describes a shell model rather than 3-D Navier–Stokes. A further fact bounds what the shell model can deliver even in principle.

The dyadic dissipation term is $\nu k_n^{2\alpha}$; the Navier–Stokes viscous term is $\nu \Delta u$, whose Fourier symbol is $-\nu|k|^2$. Hence $\alpha = 1$ *is* the Navier–Stokes dissipation — the same exponent, not an analogy. Cheskidov proves global regularity for $\alpha \geq 1/2$. Therefore:

The dyadic shell model at Navier–Stokes dissipation is provably regular.

The model’s difficult regime, $\alpha < 1/2$, corresponds to dissipation *weaker* than Navier–Stokes. A result in the open band $[1/3, 1/2)$ is a result about *hypo-dissipative* shell models. That is live mathematics — the band is open, and Cheskidov singles out $\alpha = 1/3$ as carrying the same nonlinear estimates as 4-D Navier–Stokes — but it is **not a step toward 3-D Navier–Stokes and must never be reported as one**.

An unplanned consistency check fell out of the same calculation. The K41 balance $\nu k^{2\alpha} \sim k a_k$ with $a_k \sim k^{-1/3}$ gives a dissipation scale $k_d \sim \nu^{-1/(2\alpha-2/3)}$, whose exponent *diverges* as $\alpha \rightarrow 1/3^+$: below that threshold no dissipation scale exists, so the cascade is stopped at no finite shell. That is exactly Cheskidov’s proven blow-up threshold, reached from an elementary heuristic that knows nothing of his proof.

Owner decision: this gap is to be treated as a gap to *close*, not a ceiling to accept. The programme will pursue the bridge — the $\mathbb{Z}^3$ embedding and the instantiation of the nonlinearity — both of which remain blocked on definitions that do not exist, and so admit only authored candidates submitted for human audit.

9.2 Repairing the reduction: giving the hypotheses content (Tier C)

Audit verdict B1 found the reduction’s two hypothesis parameters to be bare $\text{Prop} \rightarrow \text{Prop}$ arrows: “the Lean kernel is merely verifying a logical tautology”. The repair, carried out on 14 August 2026, is *not* to prove those steps — parking undischarged analysis in a named hypothesis parameter rather than an axiom is the pattern this programme deliberately adopted, so that the kernel tracks the debt in the type. The repair is to make the types say something.

What a compactness statement has to deliver

`IsGalerkinLimit` now carries four clauses, each load-bearing: the limit solves the untruncated system; it has *finite enstrophy* (a genuine ℓ^2 condition, since there is no cutoff to truncate the sum); that enstrophy is bounded on every horizon, with the constant quantified *inside* the horizon — compactness gives no uniformity in T , and writing $\exists C \forall T$ would claim more than the analysis provides; and the limit is genuinely the modewise limit of the Galerkin family **along a subsequence**. The subsequence is not a technicality: a compactness argument extracts one, and asserting convergence of the full sequence would overstate the mathematics.

The ℓ^2 upgrade was deliberately *localised*. Hypothesis U quantifies over *Galerkin* solutions, whose truncation clause forces $u_n = 0$ for $n > N$ — so the finite sum over $n \leq N$ already captures every nonzero mode and is *exact* there, not an approximation. Only the objects with no cutoff needed the infinite-series treatment. A bridge lemma (bounded partial sums of nonnegative terms $\iff$ summable with $\sum' \leq C$) makes the upgrade a conservative extension: nothing previously proven is invalidated.

The file remains Tier C, and two things remain unmachine-checkable

The repair makes the undischarged debt *legible*; it does not pay it. Both hypothesis parameters are still assumed. Moreover, whether the four clauses *say the right thing* is a statement-adequacy question reserved to human audit — it cannot be settled by the kernel, and that cost is scheduled rather than discovered later.

9.3 Eliminating a duplication hazard at the root

Until this point every file in `lean_src/` re-declared, verbatim, the definitions it needed from its neighbours, because the verification gate compiled each file in isolation and cross-file imports did not resolve. That convention carried a silent-drift hazard of exactly the kind this programme exists to eliminate: **a correction applied to one copy and not the other leaves both files compiling cleanly, both gates green, and the mathematics divergent**, with no mechanism able to see it.

With the project now self-contained, the gate builds the dependency graph first and then re-elaborates each file with the project’s own build directory on the search path. Imports resolve; the duplication is gone; there is exactly one definition of each object in the repository. One implementation detail is worth recording because getting it wrong would be invisible: the gate must re-elaborate with `lean` rather than rely on `lake build` alone, because a *cached* build target emits no axiom output at all — which would make the footprint check silently vacuous rather than failing loudly.

10 What Has Failed, and What That Taught

A programme that reports only successes is not reporting. The following incidents are recorded in `LL.md` with evidence, and each produced a standing process rule.

ID	Incident	Resulting rule
The largest failure to date		
—	An external expert audit killed this programme’s central framing claim (2026-08-13, verdict D1). For months the specification, this report, and the audit packet all asserted that the formalisation “restates the open core of the Millennium problem”. It did not: a 1-D index cutoff, the weight 4^n, and an unconstrained B describe a dyadic shell hierarchy. Four further verdicts (A2, A3, A5, C1) were NO, one (B1) forced the demotion of a whole file from Tier A to Tier C, and one (B2) required a structural repair. None of this was caught by two machine gates, five governing documents, or an adversarial self-review — because none of them can check whether a formal statement refers to the object it names.	Only human audit certifies statement adequacy, and it must be sought <i>early and from outside</i>. A self-review that finds and fixes real gaps (Q3–Q5) creates a false sense of completeness: the gaps it can see are not the structural ones. Corollary now in force: a file that is DRAFT pending audit may not have its informal gloss repeated in prose as though the audit had passed.

Earlier incidents

ID	Incident	Resulting rule
LL-1	A broad <code>git add -A</code> swept an agent’s mid-edit Lean file into a commit; the committed version carried <code>sorryAx</code> in five theorems. Happened <i>twice</i> in one session.	Stage by explicit path only while any background process may be writing; re-run the gate on the exact file immediately before staging.
LL-2	A submitted proposal self-reported “proven, no <code>sorry</code> ”; it did not compile (four Lean idiom errors) and its footprint contained <code>sorryAx</code> .	Never trust a self-report. Re-run the compiler on the exact artifact. Archive proposals verbatim with their compiled kernel output.
LL-3	D5’s redesign fixed stiffness, but the actual obstruction was representation-size growth.	Diagnose the failure mode before redesigning; a fix for the wrong axis is no fix.
LL-4	Two successive designs by the same author failed their feasibility premises.	A repeated failed design is a signal to <i>escalate the methodology decision</i> , not to try a third variant alone.
LL-6	A literature claim was nearly cited from memory; the fetched abstract showed the real result was narrower (well-posedness of <i>positive</i> solutions in a particular scaling range, not unconditional regularity) [9].	Verify literature precisely before citing, even for a well-known author or result.
LL-7	A top-tier-authored design memo’s own worked example quoted the wrong summation range.	The certifying agent must trace and correct rather than silently absorb; “derive before dispatch” does not remove the need for downstream re-verification.
LL-8	The verification gate itself had a false positive. A lemma proved by purely computational means reports the <i>strictly cleaner</i> footprint [<code>propext</code>]; the gate string-matched the exact three-axiom list and rejected it as unsound.	Test a changed gate predicate against a table including cases that <i>must</i> fail, before trusting a real run; and confirm the script’s exit code , not its trailing output. The rule’s wording in <code>SPEC.md</code> was amended to match its intent (membership, not equality).
—	The Mathlib dependency was never pinned. The <code>lakefile</code> said <code>require mathlib from git</code> with no revision; a cold build would have resolved to master and verified everything against a different Mathlib than all existing claims were checked with.	Pin the verifier itself by revision, and track the lockfile. An unpinned prover is a defect in the foundation of a machine-checked programme, not a build-script nuisance.
—	A proposed correction to the spec would itself have been an error (the T2 regularity class, §7). The evidence for it was a paraphrased summary, not a primary quotation.	When a correction rests on a summary rather than a direct quote, flag and re-fetch — do not edit the normative rulebook. LL-6 applied to its own application.

11 Summary of Verified State

Artifact	Content	Tier	The
CallensDualScale.lean	T-dual R_{eff} laws, Sym^2 lock, sharpness	Tier A	
DyadicShells.lean	Energy-flux telescoping, monotonicity	Tier A	
DyadicShell_Statements.lean	Shell-model Hypothesis U; concrete <code>shellB</code> and its energy conservation (DRAFT)	Tier A	
EnstrophyProduction.lean	Exact enstrophy-production identity	Tier A	
EnstrophyProductionBound.lean	Local bound $S_N^2 \leq 2\Omega_N^3$	Tier A	
AbstractAlgebraicConservation.lean	Abstract triad conservation + transversality	Tier A	
Tier A subtotal			
MillenniumReduction.lean	Conditional reduction skeleton (DEMOTED 2026-08-13, verdict B1)	Tier C	
Total kernel-compiled, no axiom outside {propext, Classical.choice, Quot.sound}			

All 72 theorems compile with clean axiom footprints; the demoted 7 remain wired into Gate 2 so that they cannot silently rot, but they carry **Tier C** claims and may not be cited as established. The two new rows relative to the previous report are `DyadicShell_Statements.lean`’s concrete nonlinearity block (`shellB`, `shellB_zero`, `sum_mul_shellB`, `shellB_energy_conservation`, `galerkin_shellB`, and the headline definition `DyadicShellHypothesisU`). Six Tier B harnesses are wired into Gate 1, each with a negative control demonstrated to fail. All gates pass at commit `a35cb92+`, re-run independently for this report.

11.1 What is genuinely established

Exact algebraic identities and inequalities for the truncated dyadic shell model — energy-flux telescoping and monotonicity, the exact enstrophy-production identity, the local production bound, and now the exact energy conservation of the concrete Katz–Pavlović nonlinearity; the T-dual effective radius laws with sharpness; a non-vacuous statement-level formalisation of Hypothesis U *for the shell model*, with nothing abstract remaining in the headline statement, and including a machine-checked refutation of a prior, provably false formulation; and exact-arithmetic verification of two structural properties of the Fourier–Galerkin nonlinearity, one of which is now also proven abstractly.

11.2 What is emphatically not established

Hypothesis U in any form — neither for Navier–Stokes nor for the shell model; any statement about actual Navier–Stokes solutions; any claim that the formalisation reduces, restates, or bears on the Millennium problem (retracted, §2); any claim that the Symmetric-Square Lock is relevant to the NSE cascade; the four tracks (all blocked on undefined objects); any uniformity verdict; the conditional Millennium reduction, whose file is now **Tier C**; and the statement-adequacy of either DRAFT file. The instantiation of B by the *true 3-D* nonlinearity — the 3-D bridge — has not been done and is open.

Definition of done, and what is absent from it

The campaign is complete when the formal skeleton is merged and human-audited, the dyadic results are kernel-verified, the datasets carry a human-issued verdict *whatever that verdict is*, the open problems are either audited or recorded as open with dated escalations, and every ledger row maps to a passing artifact. Note what is *absent*: proving Hypothesis U. If this programme produces a negative dyadic verdict or kills tracks, it **still counts as done and successful as science**. That is what falsifiable structure means — and the audit of 13 August 2026 is the first time that clause has had to be honoured against the programme’s own headline framing rather than against one of its subordinate conjectures. It was honoured: the claim was retracted the day it was received, the affected file was demoted, and the programme was re-targeted onto an object its formal artifacts actually describe.

Note on reference verification

Consistent with this programme’s own rule LL-6 (“*verify literature precisely before citing, even a well-known author or result*”), every reference below was verified by retrieving the actual arXiv abstract page, publisher listing, or official PDF at the time of writing. **No citation here is written from memory.** Where a detail could not be confirmed by direct retrieval, it is explicitly marked [UNVERIFIED] rather than being asserted. Readers should treat those specific marked details as requiring independent confirmation before reuse.

References

- [1] J. Leray. Sur le mouvement d’un liquide visqueux emplissant l’espace. *Acta Mathematica* **63** (1934), 193–248. DOI 10.1007/BF02547354. *Verified* via publisher listing and corroborated by a fetched modern survey (Ożański–Pooley, arXiv:1708.09787), which confirms Leray established local existence of strong solutions, global existence of weak solutions, weak–strong uniqueness, and a singular-set dimension bound. [UNVERIFIED]: the specific attribution of the *mollification* method was not confirmed by direct quotation of the original.
- [2] C. L. Fefferman. Existence and smoothness of the Navier–Stokes equation. Official Clay Mathematics Institute Millennium Prize problem description; also in *The Millennium Prize Problems*, CMI/AMS, 2006. *Verified* by fetching the official PDF. The problem asks for a proof of exactly one of four statements: (A) existence and smoothness on $\mathbb{R}^3$; (B) existence and smoothness on $\mathbb{R}^3/\mathbb{Z}^3$ (periodic); (C) breakdown on $\mathbb{R}^3$; (D) breakdown on $\mathbb{R}^3/\mathbb{Z}^3$. <https://www.claymath.org/wp-content/uploads/2022/06/navierstokes.pdf>
- [3] G. Prodi (1959); J. Serrin (1962). The Prodi–Serrin regularity criterion. [UNVERIFIED BY DIRECT FETCH]: the precise inequality, endpoint cases, and constants could not be confirmed from a retrievable full-text source. The criterion as used in this report — a Leray–Hopf weak solution in $L_t^s L_x^q$ with $2/s + 3/q \leq 1$ is smooth — is stated on the authority of secondary sources only. A commonly cited modern treatment is J. C. Robinson, J. L. Rodrigo, W. Sadowski, *The Three-Dimensional Navier–Stokes Equations*, Cambridge University Press, 2016 (*not independently fetched*).
- [4] T. Tao. Finite time blowup for an averaged three-dimensional Navier–Stokes equation. arXiv:1402.0290 (submitted 2014; v3 2015). *Verified* by fetching the arXiv abstract page. Constructs an averaged bilinear operator retaining the energy-identity cancellation, and exhibits finite-time blowup for the resulting *averaged* system, via an ODE tied to the Katz–Pavlović dyadic model. **It does not show blowup for the true Navier–Stokes equa-**

tions; its force is that regularity proofs must use structure finer than energy methods. [UNVERIFIED]: journal reference not listed on the abstract page.

- [5] L. Caffarelli, R. Kohn, L. Nirenberg. Partial regularity of suitable weak solutions of the Navier–Stokes equations. *Comm. Pure Appl. Math.* **35** (1982), 771–831. Pre-dates arXiv. *Verified* via journal citation and corroborating lecture notes: the one-dimensional parabolic Hausdorff measure of the singular set of any suitable weak solution is zero.
- [6] T. Buckmaster, V. Vicol. Nonuniqueness of weak solutions to the Navier–Stokes equation. arXiv:1709.10033 (2017); accepted by *Annals of Mathematics*. *Verified* by fetching the arXiv abstract page: weak solutions of 3D Navier–Stokes “are not unique in the class of weak solutions with finite kinetic energy”, proved via convex integration. [UNVERIFIED]: exact volume/page numbers.
- [7] N. H. Katz, N. Pavlović. Finite time blow-up for a dyadic model of the Euler equations. *Trans. Amer. Math. Soc.* **357**(2) (2005), 695–708; DOI S0002-9947(04)03532-9. *Verified* by fetching and reading the published PDF. Proves finite-time blow-up for the dyadic Euler equations, and for dyadic Navier–Stokes with hyper-dissipation when the dissipation degree is sufficiently small. [UNVERIFIED]: no arXiv posting could be located.
- [8] A. Cheskidov. Blow-up in finite time for the dyadic model of the Navier–Stokes equations. arXiv:math/0601074 (2006); *Trans. Amer. Math. Soc.* **360** (2008), 5101–5120. *Verified* by fetching the arXiv abstract page: finite-time blow-up for dissipation degree $\alpha < 1/3$ (improving Katz–Pavlović’s $\alpha < 1/4$), and global regularity for $\alpha \geq 1/2$.
- [9] D. Barbato, F. Morandin, M. Romito. Smooth solutions for the dyadic model. arXiv:1007.3401 (2010); *Nonlinearity*, DOI 10.1088/0951-7715/24/11/004. *Verified* by fetching the arXiv abstract page. **Scope is narrower than unconditional global regularity**: well-posedness of *positive* solutions of the viscous problem in the scaling range corresponding to Navier–Stokes, and of the inviscid problem in a suitable regularity class. This reference is the subject of lesson LL-6.
- [10] J. Bourgain, C. Demeter. The proof of the ℓ^2 decoupling conjecture. arXiv:1403.5335 (2014; v3 2015); *Annals of Mathematics* **182** (2015), 351–389. *Verified* by fetching the arXiv abstract page. Proves the ℓ^2 decoupling conjecture “for compact hypersurfaces with positive definite second fundamental form and also for the cone”.
- [11] C. Mouhot, C. Villani. On Landau damping. arXiv:0904.2760 (2009); *Acta Mathematica* **207**(1) (2011), 29–201. *Verified* by fetching the arXiv abstract page, including the journal reference. **The core result requires analytic regularity** (Gevrey-class settings are a related/limiting extension, not the central stated theorem) — see the caution in §7.
- [12] F. Golse, L. Saint-Raymond. The Navier–Stokes limit of the Boltzmann equation for bounded collision kernels. *Inventiones Mathematicae* **155** (2004), 81–161. *Verified* indirectly: no arXiv posting located for the 2004 paper, but its existence and scope are confirmed by the fetched follow-up (Golse–Saint-Raymond, arXiv:0808.0039, *J. Math. Pures Appl.* **91** (2009), 508–552), whose abstract explicitly states it completes the earlier *Invent. Math.* 155, 81–161 results. Establishes convergence of renormalised Boltzmann solutions to *Leray* solutions of Navier–Stokes.
- [13] H. Duminil-Copin, V. Tassion. A new proof of the sharpness of the phase transition for Bernoulli percolation and the Ising model. arXiv:1502.03050 (2015); *Comm. Math. Phys.* **343** (2016), 725–745. *Verified* by fetching the arXiv abstract page: finiteness of susceptibility below criticality, mean-field bounds above, and subcritical exponential decay of correlations.

- [14] A. Giveon, M. Porrati, E. Rabinovici. Target space duality in string theory. arXiv:hep-th/9401139; *Physics Reports* **244** (1994), 77–202. *Verified* by fetching and reading the review PDF. Eq. (2.2.11) establishes invariance of the free-boson partition function under $R/\sqrt{\alpha'} \rightarrow \sqrt{\alpha'}/R$ with momentum–winding exchange $m \leftrightarrow n$ — a genuine theorem about the CFT spectrum. [UNVERIFIED]: the literal phrase “self-dual radius” was not found in the pages retrieved, though the fixed-point fact itself is verified. **Important:** the “minimal length” reading is a *physical interpretation* of this exact symmetry, **not a separately proved theorem** — see §3.
- [15] C. F. Doran. Picard–Fuchs uniformization: modularity of the mirror map and mirror-moonshine. arXiv:math/9812162. *Verified* by fetching: defines $\text{Sym}^n(L)$ as the operator whose solution space is spanned by products of solutions; Theorem 5.13 states that the Picard–Fuchs equation of a family of M_n -polarized K3 surfaces is the symmetric square of a second-order Fuchsian ODE, hence *third-order* — the standard fact used in §3.
- [16] L. de Moura, S. Ullrich. The Lean 4 theorem prover and programming language. *CADE-28* (2021). <https://lean-lang.org/papers/lean4.pdf> *Verified* via converging independent listings (ACM DL, lean-lang.org).
- [17] The mathlib Community. The Lean mathematical library. *CPP 2020*, DOI 10.1145/3372885.3373824. *Verified* via converging independent listings.